

Nearest Neighbors

Nearest Neighbors

Each data point we have has
been represented using features

Stock, Date, [CalendarFeatures,
Momentum, Jump, Features of
Related stocks], Label


Nearest Neighbors

We can imagine that these features together, represent a point in an N-dimensional hypercube

Stock, Date, [CalendarFeatures,
Momentum, Jump, Features of
Related stocks], Label

N-Dimensional Hypercube

A LINE IS A 1
DIMENSIONAL
SHAPE




Any point on a line
can be represented
using 1 number

N-Dimensional Hypercube

A SQUARE IS A 2
DIMENSIONAL
SHAPE


Any point in a square
can be represented
using 2 numbers



N-Dimensional Hypercube

A CUBE IS A 3
DIMENSIONAL
SHAPE

Any point in a cube
can be represented
with 3 numbers




Nearest Neighbors

A set of N numbers represents a point in an **N -Dimensional Hypercube**


This is just to say that any data in the world can be represented as points in an N -Dimensional space

Nearest Neighbors




**All the past data
can be represented
as points in this N-
Dimensional Space**

Nearest Neighbors




When we have a new data, we can also represent that point in this space

Nearest Neighbors



The returns on the
new date can be
estimated using the
nearest neighbors

Nearest Neighbors



Predicted Return =
The weighted
average return of
the 5 nearest
neighbours

Nearest Neighbors

The trading strategy - i.e.
long or short can be decided
based on the **predicted return**

Nearest Neighbors

KNeighborsRegressor

This model in Scikit Learn can compute the **average of the labels** using the specified number of nearest neighbors

Nearest Neighbors

KNeighborsRegressor

This model in Scikit Learn can compute the **average of the labels** using the specified number of nearest neighbors

Here the average of the labels is nothing but the average of returns of nearest dates

Nearest Neighbors

KNeighborsRegressor

This model in Scikit Learn can compute the average of the labels using the specified number of **nearest neighbors**

The nearest neighbors are the dates in the past most similar to our test date

Nearest Neighbors

KNeighborsRegressor

This model in Scikit Learn can compute the average of the labels using the specified number of **nearest neighbors**

The nearer a past date is, the more similar the features of that date are to our test date

Gradient Boosted Trees

Gradient Boosted Trees

Just like **Random Forests**,
Gradient Boosted Trees
method uses an **ensemble**
of Decision Trees

Gradient Boosted Trees

There are 2 techniques involved here

Ensemble learning with boosting

Gradient Descent

Gradient Boosted Trees

There are 2 techniques involved here

Ensemble learning with boosting

Gradient Descent

RECAP

ENSEMBLE LEARNING

**INVOLVES THE USE OF MULTIPLE LEARNERS
AND COMBINING THEIR RESULTS**

RECAP

THE IDEA OF ENSEMBLE LEARNING IS SIMPLE..

MODELS TEND TO OVERFIT

IF YOU TRAIN MULTIPLE MODELS

THE OVERFITTING COMPONENTS OF EACH OF THE
MODELS WOULD BE DIFFERENT

WHEN YOU COMBINE THESE MODELS

THE OVERFITTING COMPONENTS OF THE MODELS
WOULD CANCEL EACH OTHER OUT

AND YOU ARE LEFT WITH THE COMPONENTS
THAT REALLY DESCRIBE YOUR DATA

RECAP

A MACHINE LEARNING ENSEMBLE IS A COLLECTION OF MODELS

THE MODELS IN THE ENSEMBLE CAN BE

BASED ON DIFFERENT TECHNIQUES

A COLLECTION WITH 1 SVM, 1 DECISION TREE, 1 NAIVE BAYES, 1 KNN

TRAINED ON DIFFERENT TRAINING SETS

A COLLECTION OF SVMs, EACH TRAINED ON A DIFFERENT TRAINING SET

USING DIFFERENT FEATURES

A COLLECTION OF DECISION TREES, EACH GIVEN A DIFFERENT SET OF FEATURES

USING DIFFERENT VALUES OF PARAMETERS

A COLLECTION OF K-NEAREST NEIGHBOURS, EACH WITH A DIFFERENT VALUE OF K

RECAP

**AN ENSEMBLE LEARNER
COMBINES THE RESULTS FROM
INDIVIDUAL MODELS**

THE FINAL RESULT CAN BE

A MAJORITY VOTE OF THE INDIVIDUAL MODELS

**AVERAGE OF THE RESULT FROM
INDIVIDUAL MODELS**

**A WEIGHTED FUNCTION OF THE
RESULT FROM INDIVIDUAL MODELS**

RECAP

BAGGING BOOSTING

ARE SPECIAL
ENSEMBLE
LEARNING
TECHNIQUES

BAGGING

(BOOTSTRAP-AGGREGATING)

IS AN ENSEMBLE LEARNING TECHNIQUE
THAT WAS DEVELOPED FOR CLASSIFICATION
PROBLEMS

EACH MODEL IN THE ENSEMBLE IS TRAINED
ON A DIFFERENT TRAINING SET

THESE TRAINING SETS ARE RANDOMLY
GENERATED FROM THE ORIGINAL TRAINING SET

FOR THE FINAL RESULT, EACH MODEL IS GIVEN AN
EQUAL WEIGHT AND A MAJORITY VOTE IS TAKEN

RECAP

BAGGING BOOSTING

ARE SPECIAL
ENSEMBLE
LEARNING
TECHNIQUES

Random Forest is an example of
a technique that uses bagging

RECAP

BAGGING BOOSTING

ARE SPECIAL
ENSEMBLE
LEARNING
TECHNIQUES

Gradient Boosted Trees use
the concept of Boosting

Gradient Boosted Trees

BOOSTING

**IS AN ALGORITHM FOR
ITERATIVELY ADDING LEARNERS
TO THE ENSEMBLE**

Gradient Boosted Trees

BOOSTING

is an algorithm for iteratively
identifying the trees in the
ensemble

Gradient Boosted Trees

BOOSTING

In each iteration a subset of the main training set is chosen for training

Gradient Boosted Trees

BOOSTING

In the first iteration, all the data points have equal probability of being chosen in the training set

Gradient Boosted Trees

BOOSTING

In the each subsequent iteration, the probabilities of being chosen are modified

Gradient Boosted Trees

BOOSTING

Points which have been
classified correctly more
often have lower weight

Gradient Boosted Trees

BOOSTING

Points which have been
misclassified start getting
higher weight

Gradient Boosted Trees

BOOSTING

Each subsequent tree that is built, is given a training set that has more of the points that are misclassified

Gradient Boosted Trees

BOOSTING

Finally, we have a collection of trees, each trained on a different training set

Gradient Boosted Trees

BOOSTING

The final classification is a
weighted vote of all the
trees

Gradient Boosted Trees

BOOSTING

The weight of a tree is
proportional to how accurate
the tree was on the training set

Gradient Boosted Trees

There are 2 techniques involved here

Ensemble learning with boosting

Gradient Descent

Gradient Boosted Trees

There are 2 techniques involved here

Ensemble learning with boosting

Gradient Descent

Gradient Boosted Trees

Gradient Descent

In addition to Boosting,
Gradient Boosted Trees also use
the concept of Gradient Descent

Gradient Boosted Trees

Gradient Descent

In each iteration of boosting

1. A misclassification rate is calculated
2. Weights of training samples are updated based on this rate

Gradient Boosted Trees

This is based on an error function

In each iteration of boosting

1. A misclassification rate is calculated

2. Weights of training samples are updated based on this rate

Gradient Boosted Trees

Gradient Descent

The error function is of the form

$$M(x) + C(x)$$

In each iteration of boosting

1. A misclassification rate is calculated

2. Weights of training samples are updated based on this rate

$$M(x) + C(x)$$

M is the misclassification
rate aka the loss

- 1. A misclassification rate is calculated**
2. Weights of training samples are updated based on this rate

Gradient Boosted Trees

$$M(x) + C(x)$$

C is the complexity of the trees being constructed

- 1. A misclassification rate is calculated**
2. Weights of training samples are updated based on this rate

Gradient Boosted Trees

$$M(x) + C(x)$$

Minimizing this function will minimize the error, at the same time penalizing models with very high complexity

1. A misclassification rate is calculated
2. Weights of training samples are updated based on this rate

Gradient Boosted Trees

$$M(x) + C(x)$$

This combination is what allows Boosting to find a good ensemble of trees without overfitting

- 1. A misclassification rate is calculated**
2. Weights of training samples are updated based on this rate

$$M(x) + C(x)$$

The weights are updated in such a way that this error function is minimized

1. A misclassification rate is calculated
2. Weights of training samples are updated based on this rate

$$M(x) + C(x)$$

The update is done using
Gradient descent, which is
an optimization technique

1. A misclassification rate is calculated
2. Weights of training samples are updated based on this rate

STOCHASTIC GRADIENT DESCENT

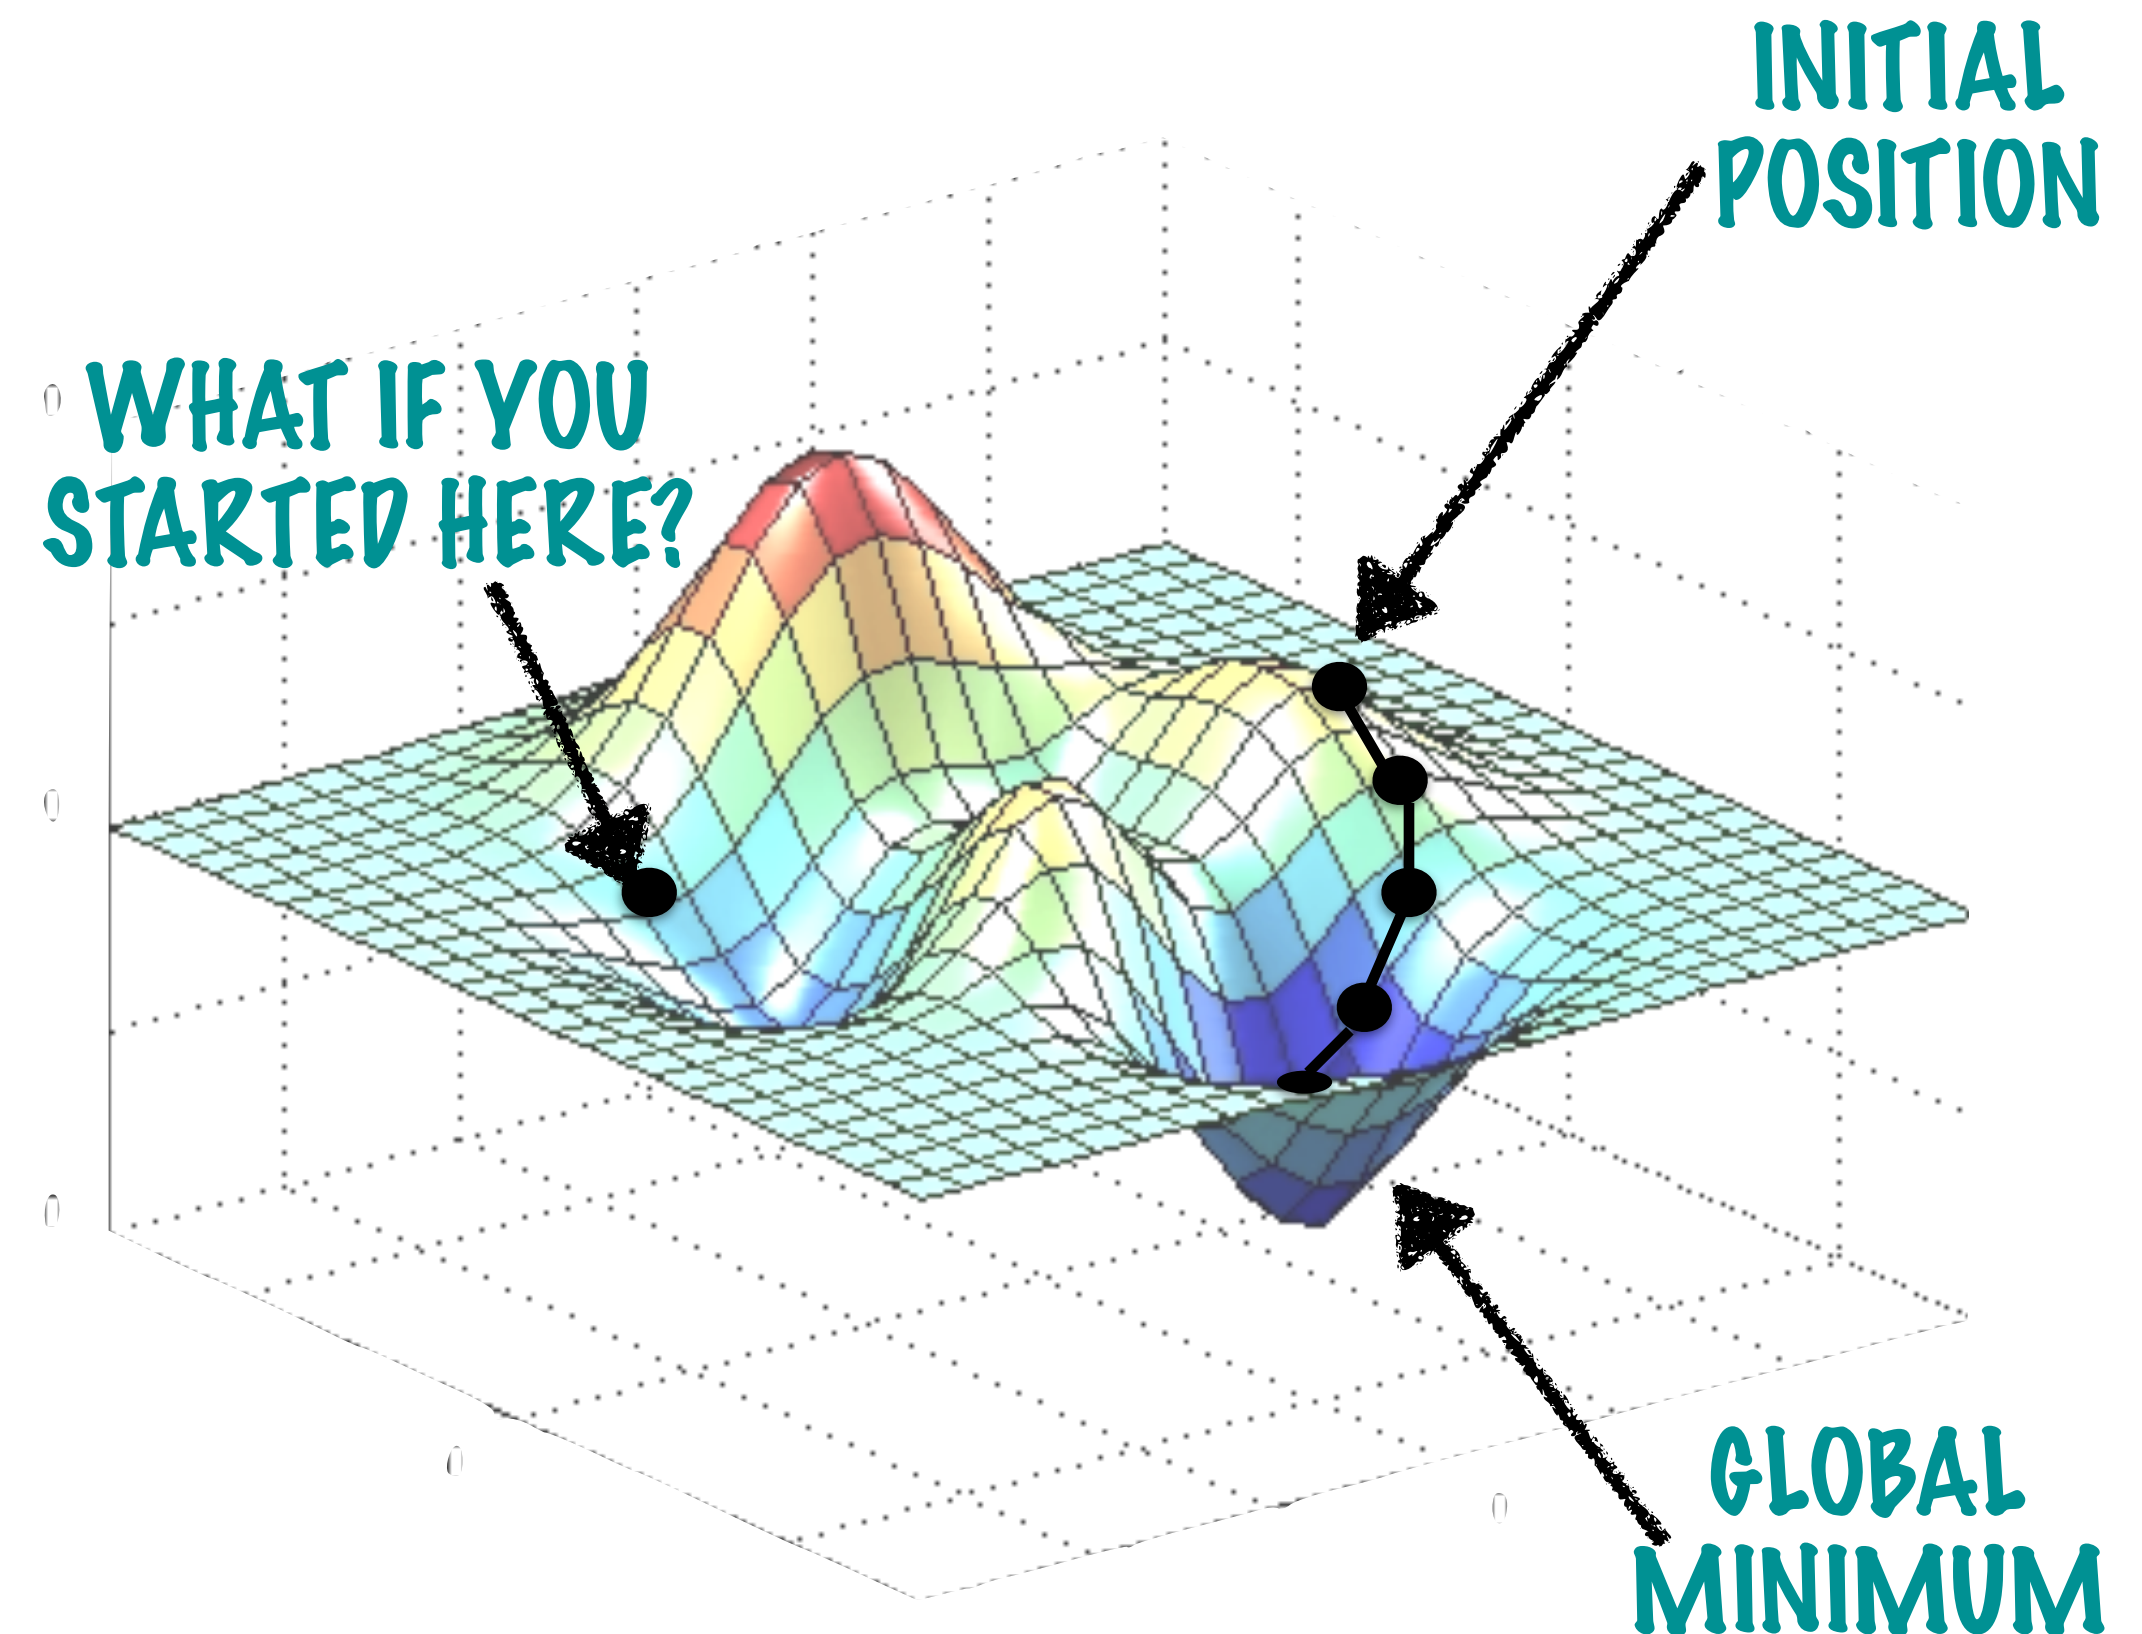
1. FIND THE CURRENT VALUE OF THE ERROR FUNCTION

2. FIND THE SLOPE AT THE CURRENT POINT AND
MOVE SLIGHTLY DOWNWARDS IN THAT DIRECTION

3. REPEAT UNTIL YOU REACH A MINIMUM

GRADIENT DESCENT DOESN'T
GUARANTEE THE GLOBAL MINIMUM

$$M(x) + C(x)$$



Gradient Boosted Trees

There are 2 techniques involved here

Ensemble learning with boosting

Gradient Descent

Gradient Boosted Trees

There are 2 techniques involved here

Ensemble learning with boosting

Gradient Descent

Gradient Boosted Trees

Ensemble learning with boosting

Gradient Descent

Gradient Boosting is not
limited to only decision trees

Gradient Boosted Trees

Ensemble learning with boosting

Gradient Descent

Using the above 2 techniques a Gradient Boosted Classifier / Regressor can be constructed with any base algorithm

Gradient Boosted Trees

XGBoost

is a library that can be used to
build Gradient Boosted Trees

Gradient Boosted Trees

XGBClassifier

is the classifier in this library that is very similar to the Classifiers in ScikitLearn

Gradient Boosted Trees

XGBClassifier

This classifier also has fit and predict methods to train and test a classifier

Gradient Boosted Trees

XGBClassifier

There are several parameters in XGBClassifier that need to be carefully chosen

Gradient Boosted Trees

XGBClassifier

n_estimators

min_child_weight

learning_rate

subsample

max_depth

gamma

colsample_bytree

Gradient Boosted Trees

XGBClassifier

n_estimators

min_child_weight

learning_rate

subsample

The number of trees to
be built

max_depth

gamma

colsample_bytree

Gradient Boosted Trees

XGBClassifier

n_estimators

min_child_weight

learning_rate

subsample

max_depth

gamma

This should depend on the
size of the training set

colsample_bytree

Gradient Boosted Trees

XGBClassifier

n_estimators

The larger the training set,
the higher the **n_estimators**
can be

Gradient Boosted Trees

XGBClassifier

n_estimators

learning_rate

max_depth

colsample_bytree

min_child_weight

subsample

After each
iteration, weights
for the training
samples are
updated

Gradient Boosted Trees

XGBClassifier

n_estimators

learning_rate

max_depth

colsample

min_child_weight

The learning
rate is a factor
that the weights
are multiplied by

Gradient Boosted Trees

XGBClassifier

n_estimators

learning_rate

max_depth

colsample_bytree

This shrinks
the weights on
each step

Gradient Boosted Trees

XGBClassifier

n_estimators

learning_rate

max_depth

colsample_bytree

min_child_weight
Such a factor can
reduce the
possibility of
overfitting

Gradient Boosted Trees

XGBClassifier

n_estimators

learning_rate

max_depth

colsample

min_child_weight

Typically, the
final values to

used are in the

range of 0.01-0.2

Gradient Boosted Trees

XGBClassifier

The maximum depth of a
tree

max_depth

Gradient Boosted Trees

XGBClassifier

This factor is used to
control over-fitting

max_depth

Gradient Boosted Trees

XGBClassifier

A higher depth means
more complex trees

max_depth

colsample_bytree

Gradient Boosted Trees

XGBClassifier

The higher the complexity, more
the chance to learn relations very
specific to a particular sample

max_depth

Gradient Boosted Trees

XGBClassifier

Typically, the value
chosen is between 3-10

max_depth

colsample_bytree

Gradient Boosted Trees

XGBClassifier

`min_child_weight`

This is another
factor that defines
the complexity of
the tree

Gradient Boosted Trees

XGBClassifier

In the tree a new child node is added to denote a new condition that has some importance

`min_child_weight`

`subsample`

`gamma`

`colsample_bytree`

Gradient Boosted Trees

XGBClassifier

The child is created by
looking at all samples
in the training set
that satisfy that
condition

min_child_weight

subsample

gamma

max_depth

sample_bytree

Gradient Boosted Trees

XGBClassifier

min_child_weight is Minimum sum of weights of all observations required in a child

Gradient Boosted Trees

XGBClassifier

n_estimators

min_child_weight

learning_rate
Higher values

subsample

max_depth
prevent
overfitting

gamma

colsample_bytree

Gradient Boosted Trees

XGBClassifier

However, this parameter requires careful tuning as too high a value can lead to underfitting

`min_child_weight`

`subsample`

`gamma`

`colsample_bytree`

Gradient Boosted Trees

XGBClassifier

This denotes the
fraction of the
training set that
should be used to
construct each tree

subsample

Gradient Boosted Trees

XGBClassifier

The lower the
value, the more
distinct each tree's
training set will be

subsample

Gradient Boosted Trees

XGBClassifier

n_estimators

min_child_weight

This is good to

subsample

max_depth

gamma

avoid
overfitting

sample_bytree

Gradient Boosted Trees

XGBClassifier

However if your
training set is too
small, then a low value
can lead to
underfitting

subsample

Gradient Boosted Trees

XGBClassifier

n_estimators

min_child_weight

Typical values

subsample

chosen are

gamma

between 0.5-1

colsample_bytree

Gradient Boosted Trees

XGBClassifier

A node in a tree is split
to make new nodes
only if it means the
error function will be
reduced by that split

gamma

Gradient Boosted Trees

XGBClassifier

Gamma specifies
the minimum loss
reduction required
to make a split

gamma

Gradient Boosted Trees

XGBClassifier

Each tree in the
ensemble can be built
using a random subset
of the overall features

`colsample_bytree`

Gradient Boosted Trees

XGBClassifier

This parameter is the fraction of the features that should be chosen to train each tree

`colsample_bytree`

Gradient Boosted Trees

XGBClassifier

Typically, a value
between 0.5-1 is
chosen

colsample_bytree

Gradient Boosted Trees

XGBClassifier

If 0.5 is chosen, then each tree is built using only 50% of the features, and the features are randomly chosen

`colsample_bytree`

Gradient Boosted Trees

XGBClassifier

The optimum value for each of these parameters has to be chosen by searching over a large number of possible combinations

Gradient Boosted Trees

Hyperopt

Hyperopt is a python library that can do an optimized search over a large number of combinations to minimize an objective

Gradient Boosted Trees

Hyperopt

```
space = {  
    'n_estimators': hp.quniform('n_estimators', 100, 1000, 1),  
    'learning_rate': hp.quniform('learning_rate', 0.025, 0.5, 0.025),  
    'max_depth': hp.quniform('max_depth', 1, 13, 1),  
    'min_child_weight': hp.quniform('min_child_weight', 1, 6, 1),  
    'subsample': hp.quniform('subsample', 0.5, 1, 0.05),  
    'gamma': hp.quniform('gamma', 0.5, 1, 0.05),  
    'colsample_bytree': hp.quniform('colsample_bytree', 0.5, 1, 0.05),  
    'nthread': 6,  
    'silent': 1  
}
```

First we define a
search space

Gradient Boosted Trees

Hyperopt

```
space = {  
    'n_estimators': hp.quniform('n_estimators', 100, 1000, 1),  
    'learning_rate': hp.quniform('learning_rate', 0.025, 0.5, 0.025),  
    'max_depth': hp.quniform('max_depth', 1, 13, 1),  
    'min_child_weight': hp.quniform('min_child_weight', 1, 6, 1),  
    'subsample': hp.quniform('subsample', 0.5, 1, 0.05),  
    'gamma': hp.quniform('gamma', 0.5, 1, 0.05),  
    'colsample_bytree': hp.quniform('colsample_bytree', 0.5, 1, 0.05),  
    'nthread': 6,  
    'silent': 1  
}
```

For each parameter a
sequence of possible
values is specified

Gradient Boosted Trees

Hyperopt

```
space = {  
    'n_estimators': hp.quniform('n_estimators', 100, 1000, 1),  
    'learning_rate': hp.quniform('learning_rate', 0.025, 0.5, 0.025),  
    'max_depth': hp.quniform('max_depth', 1, 13, 1),  
    'min_child_weight': hp.quniform('min_child_weight', 1, 6, 1),  
    'subsample': hp.quniform('subsample', 0.5, 1, 0.05),  
    'gamma': hp.quniform('gamma', 0.5, 1, 0.05),  
    'colsample_bytree': hp.quniform('colsample_bytree', 0.5, 1, 0.05),  
    'nthread': 6,  
    'silent': 1  
}
```

For instance, the number of estimators can be chosen from a sequence [100, 101,...1000]

Gradient Boosted Trees

Hyperopt

```
space = {  
    'n_estimators': hp.quniform('n_estimators', 100, 1000, 1),  
    'learning_rate': hp.quniform('learning_rate', 0.025, 0.5, 0.025),  
    'max_depth': hp.quniform('max_depth', 1, 13, 1),  
    'min_child_weight': hp.quniform('min_child_weight', 1, 6, 1),  
    'subsample': hp.quniform('subsample', 0.5, 1, 0.05),  
    'gamma': hp.quniform('gamma', 0.5, 1, 0.05),  
    'colsample_bytree': hp.quniform('colsample_bytree', 0.5, 1, 0.05),  
    'nthread': 6,  
    'silent': 1  
}
```

**quniform is used to specify that
the value of n_estimators is
chosen at random from this
sequence**

Gradient Boosted Trees

Hyperopt

```
space = {  
    'n_estimators': hp.quniform('n_estimators', 100, 1000, 1),  
    'learning_rate': hp.quniform('learning_rate', 0.025, 0.5, 0.025),  
    'max_depth': hp.quniform('max_depth', 1, 13, 1),  
    'min_child_weight': hp.quniform('min_child_weight', 1, 6, 1),  
    'subsample': hp.quniform('subsample', 0.5, 1, 0.05),  
    'gamma': hp.quniform('gamma', 0.5, 1, 0.05),  
    'colsample_bytree': hp.quniform('colsample_bytree', 0.5, 1, 0.05),  
    'nthread': 6,  
    'silent': 1  
}
```

**All values in the sequence
have equal probability of
being chosen**

Gradient Boosted Trees

Hyperopt

```
fmin(score, space, algo=tpe.suggest,  
      trials=trials, max_evals=50)
```

fmin() function will then search
through possible combinations in
this search space until an
optimum combination is found

Gradient Boosted Trees

Hyperopt

```
fmin(score, space, algo=tpe.suggest,  
      trials=trials, max_evals=50)
```

score is user defined
function whose value
needs to be minimized

Gradient Boosted Trees

Hyperopt

```
fmin(score, space, algo=tpe.suggest,  
      trials=trials, max_evals=50)
```

trials is a variable where
the result of each trial
will be stored

Gradient Boosted Trees

Hyperopt

```
fmin(score, space, algo=tpe.suggest,  
      trials=trials, max_evals=50)
```

`max_evals` is the total number of
trials the function will perform
before deciding the best
combination